

Absolute Market Unfairness, Liquidity Premiums, and the Compression of Put–Call Parity Violations in Cryptocurrency Option Markets

Balázs Králik*, Nóra Felföldi-Szűcs†, Tamás Rutai‡, Kata Váradi§

Abstract

This paper introduces *Absolute Market Unfairness* (AMU), a model-free measure of the arbitrage profit that a market maker can extract from quoted put–call-parity deviations after contract-specific conversion and trading-cost adjustments. AMU does not claim that option prices are “wrong”; it measures how much potential profit is left on the table for whoever passively supplies liquidity at the top of the book, at the expense of the taker who crosses the market. The empirical setting is the inverse BTCUSD and ETHUSD option markets of Deribit and OKX, observed from 2020 through early 2026, where contract design, collateral denomination, and settlement conventions prevent a naive transfer of textbook parity formulas. We show how AMU is computed from synchronized option and futures quotes in directly comparable basis-point units, and we ask three questions: whether AMU differs across exchanges, whether it compressed after the 2024 institutionalization of crypto markets, and whether any post-2024 change reflects a different kind of arbitrage-enforcing participant rather than merely deeper passive quoting. We find that the answer to all three is yes. AMU is minimized at the money and rises sharply away from it; it is U-shaped in maturity, elevated both at the shortest and the longest expiries and lowest around one month; it spikes in stressed markets; and it fell markedly across the 2024 event window. A panel regression with cell fixed effects shows that the post-2024 compression of the executable tail survives controls for spreads, depth, and quote staleness, which we read as evidence that parity enforcement changed structurally rather than through scale alone. Because every executable payoff is measured net of the round-trip liquidity-provision cost, AMU is not itself a liquidity premium but the mispricing that remains *over and above* the cost of supplying liquidity—the executable unfairness left in quoted prices once fair compensation for making the market has been netted out.

Keywords: cryptocurrency options, put-call parity, liquidity premium, limits to arbitrage, market efficiency, ETF approval, MiCA

JEL classification: G12, G14, G15, G23

*Corresponding Author. Assistant Professor at Corvinus University of Budapest, Institute of Finance, Fővám square 8. Budapest, Hungary, 1093. E-mail: balazs.kralik@uni-corvinus.hu.

†Associate Professor at Corvinus University of Budapest, Institute of Finance.

‡Student at Corvinus University of Budapest.

§Associate Professor at Corvinus University of Budapest, Institute of Finance.

Contents

1	Introduction	2
2	Related literature	3
2.1	Put–call parity and crypto derivative efficiency	3
2.2	Liquidity premia and the limits of arbitrage	4
3	Methodology	4
3.1	Market maker advantage and absolute market unfairness	4
3.2	The range of arbitrage configurations	6
3.3	The no-AMU band and the shape of the MMA distribution	7
3.4	Data and sample	8
3.5	Panel construction	9
3.6	Empirical design	9
4	Descriptive statistics: the AMU distribution and cross-section	10
4.1	The shape of the MMA distribution and why it peaks below zero	10
4.2	RQ1: AMU across exchanges, moneyness, and maturity	12
5	RQ2: compression after the 2024 regime shift	15
5.1	Time-series and distributional evidence	15
5.2	The dominant regime effect	17
6	RQ3: entry of new parity-enforcing capital	17
7	Interpretation: mispricing over the liquidity-provision cost	21
8	Conclusion	21

1 Introduction

Put–call parity (PCP) is one of the most basic no-arbitrage restrictions in option pricing. In a frictionless market, a European call and put with the same strike and maturity must differ in value by the forward value of the underlying. In actual markets, parity holds only inside a band: bid–ask spreads, fees, margin requirements, collateral segmentation, stale quotes, and operational frictions all determine whether an apparent wedge can be monetized.

This paper studies that wedge in cryptocurrency option markets through a measure we call *Absolute Market Unfairness* (AMU). AMU is a model-free summary of the arbitrage profit implied by quoted option and futures prices after contract-specific conversion and explicit trading-cost adjustments. It is built for crypto derivatives, where inverse settlement, non-standard collateral conventions, and fragmented exchange design make textbook parity relations insufficient on their own. By expressing the executable wedge in basis points of comparable notional, AMU turns heterogeneous option quotes into a common market-quality measure and makes cross-exchange, cross-underlying, and cross-maturity comparisons meaningful.

The interpretation of AMU is deliberately narrow, and it is worth stating precisely what AMU is not. It is tempting to file crypto parity wedges alongside familiar mature-market

anomalies such as ADR discounts, ETF premia, or on/off-the-run Treasury spreads. Those examples are not analogues of AMU. In each of them, one observes *realized trade prices* that fail to conform to a theoretical arbitrage-free relation: the same claim trades at two different prices. AMU makes no such claim about option trade prices. It says something about *quotes*: a market can be internally “fair” for a retail liquidity taker only if no executable combination of standing quotes leaves a riskless profit after costs. Whenever, for some amount of investable capital, a market can be made which—if hit—yields a riskless profit, the quoted market is transferring too much potential profit to the market maker at the expense of the retail investor. AMU measures the size of that positive tail. It is a statement about the quality and competitiveness of quoting, not about whether option prices are fundamentally mispriced.

The time-series motivation is the institutionalization of crypto markets around 2024. The approval of U.S. spot Bitcoin exchange-traded products in January 2024, the later Ether ETF milestones, and the staged implementation of MiCA plausibly expanded the set of investors able to deploy arbitrage capital in a legally and operationally robust way (U.S. Securities and Exchange Commission 2024^{b,a}, European Securities and Markets Authority 2024). Our working hypothesis is not that traders suddenly discovered parity, but that the no-arbitrage band narrowed as custody, regulation, execution, and competition improved. The cross-sectional evidence points the same way: AMU is largest exactly in the segments where arbitrage is hardest to implement, which is what one expects of a measure of mispricing *over* the cost of liquidity provision—the residual is largest where supplying liquidity and competing the wedge away is hardest.

The empirical analysis is organized around three research questions:

1. Does AMU differ across exchanges?
2. Did AMU compress after the 2024 institutionalization regime shift?
3. Is any post-2024 change consistent with the entry of a different kind of parity-enforcing liquidity supplier, rather than a pure scale increase in passive quote depth?

We answer all three in the affirmative. The remainder of the paper proceeds as follows. Section 2 places AMU in the PCP and limits-to-arbitrage literature. Section 3 defines market maker advantage (MMA) and AMU precisely, sets out the range of arbitrage configurations, and describes the data, panel, and empirical design. Section 4 presents the descriptive statistics and answers RQ1. Sections 5 and 6 answer RQ2 and RQ3 in turn. Section 7 interprets AMU as mispricing over the liquidity-provision cost through a no-arbitrage-band model. Section 8 concludes.

2 Related literature

2.1 Put–call parity and crypto derivative efficiency

Testing put–call parity has a long tradition in equity, index, and currency option markets. Early work shows that apparent parity violations are sensitive to bid–ask spreads, execution timing, short-sale constraints, taxes, and transaction costs (Klemkosky & Resnick 1979, Nisbet 1992, Kamara & Miller 1995, El-Mekkaoui & Flood 1998, Cremers & Weinbaum 2010). Crypto option markets add complications because contracts are often inverse or quanto-style instruments, collateral and settlement currencies may differ from payoff currencies, and the underlying reference index is not always directly tradable (Alexander et al. 2023, Lucic & Sepp 2024).

Recent crypto work increasingly treats derivative efficiency as a joint market-structure problem. Alexander et al. (2024) develop arbitrage tests linking Bitcoin and Ether options with perpetual futures and find that crypto derivative markets have become more efficient over time, especially for longer-maturity options. Felföldi-Szűcs (2024) study high-frequency put–call parity in Binance Ether options and emphasize that violations become much smaller once synchronization and trading costs are handled carefully. We extend this line to inverse options and matching futures on Deribit and OKX, and add a directly comparable measure of the investor-facing arbitrage profit expressed in basis points.

2.2 *Liquidity premia and the limits of arbitrage*

Reading AMU as mispricing over the liquidity-provision cost rests on the idea that arbitrage is costly and capital constrained, so that a no-arbitrage band opens up and only wedges piercing it are exploitable. Amihud & Mendelson (1986) show that bid–ask spreads affect asset pricing because investors require compensation for illiquidity. Shleifer & Vishny (1997) emphasize that textbook arbitrage is typically neither riskless nor capital-free in practice. Gârleanu & Pedersen (2011) formalize how margin requirements and funding constraints generate deviations from the law of one price. These mechanisms are directly relevant to crypto derivatives, where collateral denomination, margin rules, exchange credit risk, and stablecoin conversion all raise the cost of enforcing parity.

Two mature-market literatures show that no-arbitrage bands can persist and then move with market structure, even where AMU’s quote-side object differs from their trade-side object. Post-crisis bank balance-sheet constraints and dollar funding pressures generated persistent covered interest parity deviations (Du et al. 2018); and ETF premia and discounts can persist when the underlying is illiquid or authorized participants face balance-sheet constraints (Pan & Zeng 2017, Todorov 2021). The crypto case is, if anything, the mirror image of the CIP story: regulatory clarity and institutional access may have *reduced* the effective cost of arbitrage capital, compressing the band rather than widening it.

3 Methodology

3.1 *Market maker advantage and absolute market unfairness*

The right no-arbitrage benchmark in crypto options is still put–call parity, but its implementation is contract-specific. The most liquid contracts in our sample are inverse BTCUSD and ETHUSD options on Deribit and OKX: they are quoted against a USD index, but collateral, premium, and settlement are denominated in the underlying cryptocurrency. Parity must therefore be written after converting option premia and the futures leg into a common currency using the platform-specific inverse payoff convention (Alexander et al. 2023, Lucic & Sepp 2024). For an inverse contract this simply means that the coin-denominated premia are scaled by the conversion price S at which the coin payoff is turned into USD or a stablecoin; contract size cancels once everything is expressed in basis points of notional.

For a matched strike–expiry option pair the executable profit depends on *how* the arbitrage is implemented: forward versus backward, and whether the maker joins the call side or the put side of the book. This gives four executable paths, whose profit expressions for linear and inverse contracts are collected in Table 1. Each expression is net of a proportional round-trip trading cost and is converted to basis points by multiplying by $10,000/\text{index_level}$.

Let $m_{i,p,t}$ denote the signed basis-point payoff of path p —one of the four paths of Table 1—for option pair i at time t , after conversion and cost adjustment. We call each such (pair, path) observation a *tickpath*: every synchronized quote (“tick”) for a matched pair contributes four tickpaths, one per path, so a cell (g, t) with $n_{g,t}$ matched pairs contains $4n_{g,t}$ tickpaths. Both objects we define below are averages of m over tickpaths; they differ only in which tickpaths enter the average.

Market maker advantage (MMA) is the mean of $m_{i,p,t}$ taken over *all* tickpaths in a cell (g, t) :

$$\text{MMA}_{g,t} = \frac{1}{4n_{g,t}} \sum_{i \in g,t} \sum_{p=1}^4 m_{i,p,t}, \quad (1)$$

where $n_{g,t}$ is the number of matched pairs in the cell. Because the average includes the many paths that do not allow positive arbitrage profit to the maker after costs (but will still represent compensation for liquidity provision!), *mean MMA is negative*: it sits close to minus the round-trip cost. MMA is the average round-trip outcome of a maker who indiscriminately joins every quote and then attempts the arbitrage.

Absolute market unfairness (AMU) is the mean of $m_{i,p,t}$ taken over only the *positive-MMA tickpaths*—those with $m_{i,p,t} > 0$:

$$\text{AMU}_{g,t} = \frac{\sum_{i \in g,t} \sum_{p=1}^4 \max(0, m_{i,p,t})}{\sum_{i \in g,t} \sum_{p=1}^4 \mathbf{1}\{m_{i,p,t} > 0\}}. \quad (2)$$

AMU is the profit of a market maker who selectively joins the market only when the executable path is favorable, and is then hit. It is important to be precise about what this does and does not guarantee. AMU is *not* a promise of a given dollar profit. The realized dollar amount depends on the liquidity available along the arbitrage path and is conditional on a taker actually crossing the market. What AMU provides is information about how bad the quoted prices are: it is the average size of the exploitable, positive tail of executable payoffs left standing in the quotes. Equivalently, MMA and AMU are two moments of the same tickpath-level distribution of m : MMA is its mean over all tickpaths, AMU its mean over the positive tickpaths only.

Table 1: Executable arbitrage-profit formulas by path and contract style. To obtain the basis-point statistics used for MMA/AMU, multiply each expression by 10,000/index_level. Each expression includes the trading-cost term $\text{cost} = \text{cost_per_notional}(0.003) \times \text{index_level}$ from Table 3. Contract size cancels in the basis-point expression. For inverse contracts, S_{bid} and S_{ask} are the prices at which the coin payoff is converted to USD or stablecoin.

Path	Linear contracts	Inverse contracts
fwd, join call	$C_{bid} - P_{bid} - (F_{bid} - K)e^{-rT} - \text{cost}$	$(C_{bid} - P_{bid})S_{bid} - (F_{bid} - K)e^{-rT} - \text{cost}$
fwd, join put	$C_{ask} - P_{ask} - (F_{bid} - K)e^{-rT} - \text{cost}$	$(C_{ask} - P_{ask})S_{ask} - (F_{bid} - K)e^{-rT} - \text{cost}$
bck, join call	$P_{ask} - C_{ask} + (F_{ask} - K)e^{-rT} - \text{cost}$	$(P_{ask} - C_{ask})S_{ask} + (F_{ask} - K)e^{-rT} - \text{cost}$
bck, join put	$P_{bid} - C_{bid} + (F_{ask} - K)e^{-rT} - \text{cost}$	$(P_{bid} - C_{bid})S_{bid} + (F_{ask} - K)e^{-rT} - \text{cost}$

The trading costs are knowable but specified in a complicated, venue-dependent way. Throughout we assume a fixed proportional cost of 0.003 for a full arbitrage round trip. This is reasonable but not exact. Two properties make the choice innocuous for our purposes. First, the time variation in AMU is robust to the exact cost level: shifting the

cost floor moves all observations together and leaves the dynamics essentially unchanged. Second, holding the cost structure fixed across venues makes cross-exchange efficiency comparisons internally coherent. The cost assumption therefore serves only to place AMU on a sensible scale; none of the comparisons in this paper depend on it.

3.2 The range of arbitrage configurations

It is useful to see what AMU is measuring geometrically. For a given option leg—say the call—the quotes define a *real* market $[C_{bid}, C_{ask}]$, while parity combined with the put and the matching futures defines a *synthetic* market $[C_{bid}^{syn}, C_{ask}^{syn}]$, where $C^{syn} = P + (F - K)e^{-rT}$. Both intervals live on the same price axis, and their relative position determines whether, and how strongly, an arbitrage exists. Figure 1 shows the three qualitatively distinct configurations.




Figure 1: Real quoted versus parity-implied (synthetic) option market. When the markets are crossed and disjoint ($C_{ask} < C_{bid}^{syn}$, or the symmetric case), there is a single guaranteed riskless arbitrage. When one market is nested inside the other, or when the two straddle each other, there is no guaranteed profit but two paths carry a *potential* market-maker advantage, realized only if the maker's quote is joined and then hit. In the nested case these are the forward and backward join-call paths ($C_{bid}^{syn} - C_{bid}$ and $C_{ask} - C_{ask}^{syn}$), with the join-put payoffs equal to their negatives; in the straddling case they are the forward join-call and forward join-put paths ($C_{bid}^{syn} - C_{bid}$ and $C_{ask}^{syn} - C_{ask}$). The latter two configurations are what AMU aggregates. Schematic, illustrative levels.

In the crossed, disjoint case the real ask lies below the synthetic bid (or vice versa); a trader can simultaneously buy in one market and sell in the other and lock in the gap. This is the strongest, guaranteed form of arbitrage. In the nested case one interval contains the other, and in the straddling case the two intervals overlap only partially. In both of the latter two configurations there is no guaranteed profit: whether the maker earns depends on which side is joined and whether a taker crosses. These are exactly the states that carry a positive but conditional payoff, and they are what MMA and AMU aggregate over the whole surface.

It is important to emphasize that MMA (and AMU) do not measure the *width* of the bid-ask spread, i.e. the compensation market makers demand for liquidity provision. Rather, MMA is a measure of how *different* this compensation is between the real and synthetic markets. If they are different, the liquidity taker is potentially harmed. A good illustration for this problem is the U.S. equity National Best Bid and Offer (NBBO) rule under Regulation NMS, which prohibits trading through a better-priced quote on another

venue: MMA is precisely the through-price gap that such a cross-market best-execution rule is designed to eliminate, except that here the competing “venue” is the parity-implied synthetic option market rather than a second exchange (U.S. Securities and Exchange Commission 2005).

3.3 The no-AMU band and the shape of the MMA distribution

The configurations in Figure 1 raise a sharper question: for a *given* synthetic market, how far can a real quote move before it becomes unfair? The answer pins each real quote to its synthetic counterpart within a tolerance set by the trading cost, and it fixes the shape that the MMA distribution “should” take.

Consider the real call ask C_{ask} . Two of the four paths in Table 1 contain it—forward/join-put and backward/join-call. Requiring both payoffs to be non-positive (no AMU) yields

$$P_{ask} + (F_{ask} - K)e^{-rT} - \text{cost} \leq C_{ask} \leq P_{ask} + (F_{bid} - K)e^{-rT} + \text{cost}. \quad (3)$$

Both bounds use the *same-side* put quote P_{ask} , so the band is centred on the synthetic ask evaluated at the futures mid, $C_{ask}^{\text{syn}} \equiv P_{ask} + (F_{mid} - K)e^{-rT}$, with half-width $\text{cost} - \frac{1}{2}(F_{ask} - F_{bid})e^{-rT}$. The symmetric calculation pins C_{bid} to C_{bid}^{syn} . In words: a fair real market must track the synthetic market quote by quote, to within $\pm \text{cost}$ (Figure 2).

The role of the cost is worth emphasizing, because it is easy to mis-state. It is *not* that C_{ask} merely lies somewhere inside the synthetic bid–ask; that would be a statement about the synthetic spread alone and would carry no cost at all. Rather, C_{ask} is tied to the synthetic *ask* within a $\pm \text{cost}$ cushion, and that cushion is exactly what gives the band its width. Setting $\text{cost} = 0$ makes the point vivid: whenever the futures market has a spread ($F_{ask} > F_{bid}$), the interval in (3) inverts and becomes empty, so no strictly fair ask exists without a cost cushion. As all frictions vanish together, the band collapses and the real market is forced onto the synthetic one—the law of one price.




Figure 2: No AMU requires each real quote to track its matching synthetic quote within $\pm \text{cost}$. The band centres are the frictionless (no-cost) synthetic prices $P_{bid} + (F_{bid} - K)e^{-rT}$ and $P_{ask} + (F_{ask} - K)e^{-rT}$; the trading cost buys the $\pm \text{cost}$ cushion around each. As costs shrink, the real market is squeezed onto the synthetic market. Schematic, illustrative levels.

Equation (3) also fixes the shape of the MMA distribution. Write $u = C_{ask} - C_{ask}^{\text{syn}}$; up to the futures half-spread the two conjugate C_{ask} payoffs are $u - \text{cost}$ (forward, join put) and $-\text{cost} - u$ (backward, join call). As C_{ask} ranges over its fair band $u \in (-\text{cost}, \text{cost})$, each sweeps linearly over $(-2\text{cost}, 0)$. If makers placed quotes with no preferred point inside the band, the implied MMA would be *uniform* on $(-2\text{cost}, 0)$: a rectangle of width

2 cost immediately to the left of zero. That rectangle is the no-information benchmark (dashed in Figure 3).

The observed distribution is not flat, however; it is sharply peaked at $-\text{cost}$, and a short equilibrium argument says why. The two conjugate paths share the call-ask and put-ask quotes, and their payoffs sum to -2 cost , so the two MMAs are mirror images about $-\text{cost}$: whatever advantage one side gains, the other loses. Each maker would like a higher MMA and improves their quote to get it, but that lowers the conjugate MMA by the same amount. The put-ask quoter (the forward arbitrageur) and the call-ask quoter (the backward arbitrageur) are thus quoting against each other, and the only outcome in which neither hands advantage to the other is the symmetric one, where both carry the *same* MMA. That midpoint is $-\text{cost}$. The identical argument applied to the call-bid and put-bid quotes pins their conjugate paths at $-\text{cost}$ as well. Competitive quoting therefore parks the modal MMA at $-\text{cost}$ —exactly the peak in Figures 3 and 4—rather than spreading it uniformly across the fair band. Stale or aggressive quotes that momentarily break the symmetry populate the tails, and the part of the right tail that crosses zero is the executable AMU.




Figure 3: Schematic MMA distribution. The dashed rectangle over $(-2 \text{ cost}, 0)$ is the no-information benchmark implied by the no-AMU band (uniform quoting). Competitive quoting instead concentrates the mode at $-\text{cost}$, the symmetric equilibrium of the conjugate forward/backward paths; the mass that crosses zero into positive territory is the executable AMU. Schematic, illustrative levels.

3.4 Data and sample

The empirical sample uses high-frequency Deribit and OKX BTCUSD and ETHUSD inverse options and matching expiry futures from Tardis.dev (Tardis.dev 2026), spanning 2020 through February 2026. For each exchange, underlying, strike, maturity, and timestamp, the data contain synchronized best bid and ask quotes for calls, puts, and the matching-expiry futures. Raw tick data are processed with our own open-source codebases—`balazs14/TardisProject` and `balazs14/publications` (subfolder `AMU-liquidity-premium`)—which implement synchronization, filtering, parity-path construction, and AMU aggregation (Králík 2026*b,a*). The same AMU calculation is applied consistently across the full horizon, so pre-2024 and post-2024 comparisons are internally coherent.

Table 2 summarizes coverage. Deribit lists a wider surface than OKX—on the order of 114 versus 50 contracts per day for BTCUSD and 66 versus 34 for ETHUSD—while average call and put spreads are broadly comparable across venues at roughly 56–66 basis points, slightly wider for ETH on OKX. The OKX series are longer (about 2,269 and 2,145 trading days for BTC and ETH) than the Deribit series (about 1,139 days each), reflecting the fact that deribit started offering matching expiration futures later than OKX.

	BTCUSD		ETHUSD	
	deribit	okex	deribit	okex
Avg num contracts	109	83	67	52
Avg num expirations	7	6	6	4
Avg num strikes	80	57	53	46
Avg call spread (bp)	46	35	46	36
Avg put spread (bp)	42	29	53	32
Num days in sample	366	366	366	366

Table 2: Summary statistics of matched put–call pairs in the 2020–2026 sample.

3.5 Panel construction

Observations are synchronized on a regular time grid, filtered to matched put–call pairs with the same strike and expiry and usable bid and ask quotes, and aggregated into cells indexed by (day, exchange, ref_sym, rel_strike_bucket, tte_bucket). This grouping preserves the cross-section needed to compare liquidity states while supporting a time-series analysis of the regime shift. The core data-processing assumptions are collected in Table 3.

Table 3: Data-processing assumptions for MMA/AMU construction.

Component	Assumption
Trading-cost parameter	<code>cost_per_notional</code> = 0.003. The cost term is <code>cost_per_notional</code> multiplied by index level.
AMU clipping rule	MMA series are clipped to [0, 100] basis points before positive-part averaging.
MMA clipping rule	MMA series are clipped to [−500, +500] basis points.
Spread filter	Call and put spreads are capped at 1000 basis points.
Relative-strike filter	Allowed range is [0.79, 1.21].
Minimum quote size	<code>min_quote_size_dollar</code> must be at least 5000.
Subsampling	Last-known observation carried forward to the configured interval (default 5min); a leg with a missing or non-updated bid/ask price or size is marked stale.

3.6 Empirical design

The dependent variable in the regime regressions is the cell-level positive-part mean AMU, $AMU_{g,t}$ —a quote-level, panel-observable object rather than a latent dealer-balance-sheet

quantity. The regressors are the post-2024 indicator, its segment interactions, and three liquidity-friction controls: the average put–call spread, a log quote-depth proxy, and a stale-quote fraction:

$$\text{Spr}_{g,t} = \frac{1}{2}(\text{mean_call_spread_bp}_{g,t} + \text{mean_put_spread_bp}_{g,t}), \quad (4)$$

$$\text{Depth}_{g,t} = \log(\text{mean_min_quote_size_dollar}_{g,t}), \quad (5)$$

$$\text{Stale}_{g,t} = \text{mean stale-quote fraction across call, put, and spot legs.} \quad (6)$$

It is important to be clear about what AMU is here, because a purely limits-to-arbitrage reading can slide into treating AMU as a liquidity premium—the compensation a market maker earns for supplying liquidity. It is not that. The compensation for liquidity provision is the round-trip cost term, which is already *subtracted* inside every path payoff in Table 1. AMU is the executable payoff that survives *over and above* that cost term: it is the mispricing left once fair compensation for making the market has been netted out. In the language of limits-to-arbitrage models (Amihud & Mendelson 1986, Shleifer & Vishny 1997, Gârleanu & Pedersen 2011), the cost term is the width of the no-arbitrage band that liquidity provision alone would justify, and AMU is the part of the executable wedge that pierces that band. The controls therefore test not whether AMU *is* a liquidity premium, but whether this residual mispricing rises with the frictions—wider spreads, thinner depth, staler quotes—that make it harder to compete the wedge back inside the cost band.

The institutional event dates are January 10, 2024 (U.S. spot Bitcoin ETP approval), the May/July 2024 spot Ether ETF approval and trading milestones, June 30, 2024 (MiCA stablecoin provisions), and December 30, 2024 (broader MiCA CASP framework). Consistent with the literature on institutional frictions, these are treated as regime dates rather than single-cause shocks, and are introduced through a post-2024 indicator and event-window dummies matched to panel day.

We estimate two within-cell specifications, both with cell fixed effects γ_g and HC1 standard errors. The first (Section 5) isolates the regime effect by regressing AMU on the post-2024 indicator and its segment interactions:

$$\text{AMU}_{g,t} = \beta_0 + \beta_1 \text{Post}_t + \beta_2 (\text{Post}_t \times \text{ETH}_g) + \beta_3 (\text{Post}_t \times \text{ATM}_g) + \beta_4 (\text{Post}_t \times \text{NearExp}_g) + \gamma_g + \varepsilon_{g,t}. \quad (7)$$

The second (Section 6) replaces the interactions with the liquidity-friction controls, to test whether the regime effect survives them:

$$\text{AMU}_{g,t} = \beta_0 + \beta_1 \text{Post}_t + \beta_2 \text{Spr}_{g,t} + \beta_3 \text{Depth}_{g,t} + \beta_4 \text{Stale}_{g,t} + \gamma_g + \varepsilon_{g,t}. \quad (8)$$

Because the cell fixed effects γ_g absorb all time-invariant segment heterogeneity (including the ETH, moneyness, and maturity main effects), β_1 measures a within-cell level shift. Table 4 defines the abbreviated regressor names.

4 Descriptive statistics: the AMU distribution and cross-section

4.1 The shape of the MMA distribution and why it peaks below zero

The single most informative object in the data is the shape of the MMA distribution, shown in Figure 4. Pooled across all markets and all four paths, the density is sharply peaked *at minus the trading cost*—near -30 basis points for our 0.003 assumption—and

Table 4: Abbreviations used in the empirical specifications.

Abbreviation	Definition
Post_t	Post-2024 regime indicator (1 for dates on/after 2024-01-10, 0 otherwise).
$\text{Spr}_{g,t}$	Average put-call spread in basis points, $0.5 \times (\text{mean_call_spread_bp} + \text{mean_put_spread_bp})$.
$\text{Depth}_{g,t}$	Log quote-depth proxy, $\log(\text{mean_min_quote_size_dollar})$, min taken over put/call bid/ask quote sizes.
$\text{Stale}_{g,t}$	Mean stale-quote fraction across call, put, and spot legs.
ETH_g	ETH-underlying segment indicator.
ATM_g	At-the-money indicator, $ \text{rel_strike_bucket} - 1 \leq 0.05$.
NearExp_g	Near-expiry indicator, $\text{tte_bucket} \leq 0.35$ years.
γ_g	Cell fixed effect for the tuple $(\text{exchange}, \text{ref_sym}, \text{rel_strike_bucket}, \text{tte_bucket})$.
$\varepsilon_{g,t}$	Idiosyncratic error.

falls away on both sides. The vertical line at zero separates the region where no maker arbitrage is possible (to the left, where executable payoffs are negative net of cost) from the positive tail (to the right), which is the executable region that feeds AMU.




Figure 4: Average density of the four MMA paths across all markets. The distribution peaks at minus the round-trip cost (the “no maker arbitrage” line); only the positive tail to the right of zero is executable and contributes to AMU.

This shape has a clean economic reading, and it is central to the whole paper. Market makers improve their quotes only as much as they need to, given costs. Competition in speed and inventory management pushes quotes toward the point where the expected marginal profit of quoting more tightly just meets the implementation-cost floor. That is why the mass concentrates at $-\text{cost}$: in the modal state the maker has competed the wedge down to roughly the cost of doing the trade, and no more. The distribution is not a point mass, however, because quoting is noisy. Occasionally a maker quotes a price

that is *too good* for the buy side—these observations populate the left tail, below the cost floor—and occasionally a maker quotes a price that is *too bad* for the buy side, leaving an executable wedge. Those “too bad” quotes are the right tail, and their conditional average size is AMU. The reason AMU is not zero is simply that repricing is not instantaneous: for a short interval after a shock, some quotes are stale or mispriced relative to the synthetic market. It follows that the quicker and denser the competition among makers, the faster both tails are erased, and in particular the smaller the “too bad” tail becomes. AMU is therefore a direct, sign-restricted read on how much of that transient, taker-adverse mispricing survives in a given segment or period. Everything that follows is an application of this one idea: measuring the right tail across exchanges, across the surface, and across time.

4.2 RQ1: AMU across exchanges, moneyness, and maturity

The global averages in Table 5 establish that exchanges differ, though the differences are moderate and not a simple BTC-versus-ETH split. Mean AMU is about 6.8 bp for BTCUSD on Deribit and 6.7 bp on OKX, and about 6.7 bp for ETHUSD on Deribit, but noticeably higher at 7.8 bp for ETHUSD on OKX. The standout segment is thus ETH on OKX, the thinnest of the four books (Table 2), which is what one expects if AMU is mispricing over the liquidity-provision cost: the least deep venue—underlying combination leaves the largest executable wedge above that cost. Positive MMA is far from rare: across the four books, between roughly 37% and 41% of synchronized path observations are executable before conditioning, which is why the positive-part average is an economically meaningful object rather than a measure of isolated glitches.

		mean AMU(bps)	positive MMA occurs on any path	Num Observations
underlying	exchange			
BTCUSD	deribit	31.73	4,852,286	11,541,389
	okex	33.08	3,084,146	8,765,581
ETHUSD	deribit	30.45	3,174,854	7,056,276
	okex	31.41	1,617,190	5,508,204

Table 5: Global AMU summary across exchanges and underlyings.

The cross-sectional shape of AMU over the option surface is even sharper than the venue averages, and it is highly regular across all four exchange–underlying combinations. Figure 5 shows a pronounced V-shape in relative strike: AMU is minimized at the money, around 3–4 bp, and rises steeply into the wings, reaching roughly 20–27 bp at relative strikes of 0.8 and 1.2. Figure 6 shows that AMU is U-shaped in time to expiry: it is elevated at the very shortest maturities (about 8 bp as expiry approaches), falls to a minimum of roughly 3–4 bp around one month, and then rises again for longer-dated contracts, reaching 12–18 bp at the far end. This maturity pattern is worth flagging because it corrects a natural but wrong intuition: AMU is *not* monotonically lower for longer-dated options. It is smallest for near-dated, near-the-money contracts—the deepest and most actively arbitrated part of the surface—and largest in the two illiquid corners, the shortest and the longest expiries, and the deep wings.




Figure 5: AMU by relative strike across exchanges and underlyings. AMU is minimized at the money and rises sharply into the wings.




Figure 6: AMU by time to expiry across exchanges and underlyings. AMU is U-shaped: elevated at the shortest and longest maturities, lowest around one month.

Taken together, these regularities answer RQ1 and set up the rest of the paper. AMU differs across venues, but far more of its variation is explained by *where on the surface* a contract sits—how far from the money and how far from a one-month horizon—than by which exchange lists it. Both patterns line up with liquidity: the executable wedge is largest exactly where books are thin and quotes stale.

5 RQ2: compression after the 2024 regime shift

5.1 Time-series and distributional evidence

The time-series evidence for compression is direct. Figure 7 plots daily AMU by exchange and underlying against the BTCUSD index over the full sample. Through 2020–2023, AMU is both elevated and highly volatile, with frequent spikes into the 30–50 bp range during turbulent episodes; from 2024 onward it drops to a much lower and tighter regime, mostly below 10 bp and often in the 2–5 bp range. The spikes in the earlier period cluster in stressed markets, confirming that AMU rises when market making becomes balance-sheet intensive and quote quality deteriorates.




Figure 7: Daily AMU by exchange and underlying, 2020–2026, with the BTCUSD index overlaid. AMU is high and volatile before 2024 and compresses to a lower, tighter band afterward.

Zooming in on the transition, Figure 8 tracks the all-market average across 2023–2024 with the institutional event dates marked. Average AMU runs at roughly 15–20 bp through 2023, steps down around the January 2024 Bitcoin ETP approval, and continues to fall through the Ether ETF and MiCA milestones to about 2–5 bp by late 2024. The event-study windows in Figure 9 are noisier, as one would expect from overlapping 120-day windows around five distinct dates, but they show the same tendency: AMU is generally lower and less spiky in the post-event windows than in the pre-event windows.

Two further figures show *how* the compression happened, distributionally and across the surface. Figure 10 overlays the pre- and post-2024 MMA histograms. The post-2024 density is markedly more concentrated at the cost floor, with thinner tails on both sides; the positive-tail average falls from 53.9 bp before 2024 to 40.7 bp after. This is precisely




Figure 8: All-market average AMU, 2023–2024, with the main institutional event dates marked by vertical lines.




Figure 9: Event-study windows for AMU around the main 2024 institutional event dates, shown separately by event.

the signature predicted by the mechanism of Section 4.1: faster, denser competition among makers erases more of the transient “too good” and “too bad” quotes, pulling the whole distribution toward the no-arbitrage floor. Figure 11 shows the same compression over the strike–maturity grid. Before 2024 a hot band of high AMU sits at longer maturities and away-from-money strikes; after 2024 the interior of the grid cools sharply, and only the extreme deep-out-of-the-money, long-dated corner remains elevated. The decomposition in Figure 12 confirms that the change is broad but uneven: AMU fell in every segment, with the largest decline in the far-expiry bucket and the smallest on Deribit, and somewhat larger declines for ETH than BTC and for non-ATM than ATM strikes—again, compression is largest where the pre-2024 wedge was largest.




Figure 10: Total MMA histogram across all four paths, split into pre- and post-2024 periods. The post-2024 distribution is more concentrated at the cost floor, and the positive-tail average (AMU) falls from 53.9 to 40.7 bp.

5.2 The dominant regime effect

Table 6 makes the compression precise. It regresses cell-level AMU on the post-2024 indicator and its segment interactions (Equation 7), with cell fixed effects. The Post_t coefficient is -0.509 bp ($t = -20.6$), by a wide margin the largest and most significant term—a within-cell level shift of about 9% of the standard deviation of AMU that is common to the whole surface. The interactions only modulate it. The near-expiry interaction is positive and sizeable ($+0.410$, $t = 17.7$), so the shortest-dated cells shared least in the compression, echoing the heatmap; the ETH interaction is small and negative (-0.051); and the ATM interaction is insignificant. For RQ2 the conclusion is unambiguous: the post-2024 regime is a dominant, broad-based compression of the executable wedge, not a segment-specific artifact.

6 RQ3: entry of new parity-enforcing capital

The compression documented in RQ2 could in principle be nothing more than “more of the same” liquidity: if spreads tightened, books deepened, and quotes became fresher, the




Figure 11: Mean AMU across relative-strike and maturity buckets, pre- versus post-2024. The high-AMU interior of the surface cools after 2024; only the deep-OTM, long-dated corner stays elevated.




Figure 12: Pre- to post-2024 change in AMU by segment (BTC/ETH, Deribit/OKX, ATM/Non-ATM, NearExp/FarExp). Every segment compresses; the far-expiry bucket compresses most.

Term	Coef	$\beta \text{std}(X)/\text{std}(Y)$
Post_t	-21.103*** (0.772)	-12%
$\text{Post}_t \times \text{ETH}_g$	-0.293 (0.931)	-1%
$\text{Post}_t \times \text{ATM}_g$	8.368*** (1.118)	12%
$\text{Post}_t \times \text{NearExp}_g$	-5.098*** (0.893)	-9%
Total R^2	0.264	
Observations	112,025	
Fixed effects	cell_id	

Table 6: RQ2 regression: cell-level AMU on the post-2024 indicator and its segment interactions (Equation 7), cell fixed effects, HC1 standard errors in parentheses. Stars denote p -value bands based on the t -statistic; the right column reports the standardized effect $\beta \text{std}(X)/\text{std}(Y)$.

executable tail would shrink mechanically. RQ3 asks whether anything survives once those channels are held fixed. Table 7 regresses AMU on the post-2024 indicator together with the three liquidity-friction controls (Equation 8), again with cell fixed effects.

Two things stand out. First, the friction controls enter with the expected signs and are highly significant: staleness raises AMU strongly (+1.499, $t = 13.6$), exactly as the mechanism of Section 4.1 predicts—the executable tail is made of transiently mispriced, slow-to-update quotes—while wider spreads raise it marginally and deeper books lower it. Much of the raw regime effect does run through these channels: the Post_t coefficient falls from -0.509 in Table 6 to -0.087 here once frictions are held fixed. Second, and decisively, that residual -0.087 is still negative and strongly significant ($t = -6.8$). A part of the compression is therefore *not* explained by observable liquidity improvement. Were the post-2024 change merely a scale-up of the same passive quoting, it would operate entirely through spreads, depth, and staleness, and the residual Post_t would be indistinguishable from zero. A significant residual is the signature of a structural change in *who* enforces parity and how quickly—the entry of a different class of arbitrage-active participant. Figure 13 shows how tightly the executable wedge is tied to the spread channel in the first place, which is what makes the surviving residual notable.

Putting the three questions together: exchanges differ, but surface location matters more (RQ1); AMU compressed sharply and broadly across the 2024 event window, with Post_t the dominant term (RQ2); and a statistically significant part of that compression survives the friction controls (RQ3), which we read as evidence that a different kind of parity-enforcing participant entered the market, rather than simply more of the same passive depth.

Term	Coef	$\beta \text{std}(X)/\text{std}(Y)$
Post_t	-9.126*** (0.392)	-5%
$\text{Spr}_{g,t}$	0.180*** (0.001)	55%
$\text{Depth}_{g,t}$	0.143 (0.087)	1%
$\text{Stale}_{g,t}$	7.091*** (0.938)	2%
Total R^2	0.472	
Observations	112,025	
Fixed effects	cell_id	

Table 7: RQ3 regression: cell-level AMU on the post-2024 indicator and the liquidity-friction controls (Equation 8), cell fixed effects, HC1 standard errors in parentheses.




Figure 13: Mean MMA against average put–call spread and log quote size, and by underlying and exchange. Mean MMA falls almost one-for-one with the spread and is roughly flat in depth, showing how strongly the executable wedge is tied to the friction channel that RQ3 controls for.

7 Interpretation: mispricing over the liquidity-provision cost

It is tempting to read AMU as a liquidity premium, but that conflates two distinct objects. The liquidity premium is the compensation a maker earns for supplying liquidity, and in our construction that compensation is exactly the round-trip cost term netted out of every path payoff. AMU is what is left *after* that netting: the executable, risk-free profit that a maker who passively joins the top of the book can lock in *over and above* fair compensation for making the market, because competitors do not instantly close the wedge. It is not a pricing error in the fundamental-valuation sense either. It is mispricing relative to the parity-implied synthetic market, measured net of the cost of liquidity provision—the “unfairness” in the name.

Formally, the no-arbitrage restriction is a band. For a parity wedge $X_{i,t}$,

$$|X_{i,t}| \leq B_{i,t}, \quad B_{i,t} = \tau_{i,t} + s_{i,t} + m_{i,t} + f_{i,t} + \ell_{i,t} + \kappa_{i,t}, \quad (9)$$

where τ denotes explicit fees, s bid–ask and market-impact costs, m margin and collateral costs, f funding and stablecoin-conversion costs, ℓ legal and regulatory uncertainty, and κ counterparty, custody, and operational costs. This band width *is* the liquidity-provision cost; a wedge inside it is fair compensation for making the market. AMU is the sign-restricted content that *pierces* the band: the positive executable payoff remaining once $B_{i,t}$ has been netted out, i.e. the mispricing over and above the cost of supplying liquidity. Empirically we net only the single proportional cost term of Table 1, so AMU is a conservative, lower-bound read on that residual. The claim of Section 5 is that several components of $B_{i,t}$ fell after 2024, tightening the band against which AMU is measured.

A compact way to see the mechanism is to let the band depend on arbitrage capacity,

$$B_{i,t} = \frac{\phi_{i,t}}{A_t}, \quad (10)$$

where $\phi_{i,t}$ is the per-unit implementation cost and A_t is effective arbitrage capital. Regulatory clarity and ETF-linked integration lower $\phi_{i,t}$ and raise A_t , both of which narrow the band. The same logic delivers a testable dynamic prediction: treating deviations as persistent inside the band and mean-reverting outside it,

$$X_{i,t+1} = \begin{cases} \rho_0 X_{i,t} + \varepsilon_{i,t+1}, & |X_{i,t}| \leq B_{i,t}, \\ \rho_1 X_{i,t} + \varepsilon_{i,t+1}, & |X_{i,t}| > B_{i,t}, \end{cases} \quad |\rho_1| < |\rho_0|, \quad (11)$$

a downward break in $B_{i,t}$ at the 2024 threshold implies both a leftward shift in the distribution of $|X_{i,t}|$ and faster mean reversion of large deviations. The first implication is exactly the compression documented in Figures 7–12 and in the AMU regression; the second is consistent with the concentration of the post-2024 MMA distribution at the cost floor. This is the sense in which the collapse of parity violations reflects a transition from visible-but-non-bankable arbitrage to institutionally enforceable arbitrage.

8 Conclusion

This paper introduces Absolute Market Unfairness as a model-free way to measure the executable, taker-adverse tail of put–call-parity wedges in inverse cryptocurrency option markets, after the relevant contract conversions and trading costs are imposed. AMU does

not assert that option prices are wrong; it measures how much potential profit the quoted market hands to a passive liquidity supplier at the expense of the taker who crosses it, without guaranteeing any particular dollar amount.

Using Deribit and OKX BTCUSD and ETHUSD options from 2020 through early 2026, we answer three questions. Exchanges differ, but the executable wedge is governed far more by location on the option surface—smallest for near-dated, near-the-money contracts and largest in the illiquid corners—than by venue. AMU compressed sharply across the 2024 institutionalization window, from a volatile 15–50 bp regime to a tight 2–5 bp regime, in a way that tracks the ETF and MiCA event dates. And the compression of the executable tail survives controls for spreads, depth, and staleness within panel cells and is concentrated in the least-liquid segments, which we interpret as statistically significant evidence that a different kind of parity-enforcing participant entered the market rather than simply more passive depth. The value of the AMU framework is that it moves the analysis from cataloguing parity violations to measuring, in comparable basis points, the mispricing that survives over and above the cost of supplying liquidity: how much executable unfairness remains once fair compensation for market making has been netted out, and how that residual responds to regulation, custody, capital access, and exchange competition.

References

- Alexander, C., Chen, D. & Imeraj, A. (2023), ‘Crypto quanto and inverse options’, *Mathematical Finance* **33**(4), 1005–1043.
- Alexander, C., Deng, J. & Imeraj, A. (2024), ‘Arbitrage opportunities and efficiency tests in crypto derivatives markets’, *Journal of Financial Markets* **71**, 100938.
- Amihud, Y. & Mendelson, H. (1986), ‘Asset pricing and the bid-ask spread’, *Journal of Financial Economics* **17**(2), 223–249.
- Cremers, M. & Weinbaum, D. (2010), ‘Deviations from put-call parity and stock return predictability’, *Journal of Financial and Quantitative Analysis* **45**(2), 335–367.
- Du, W., Tepper, A. & Verdelhan, A. (2018), ‘Deviations from covered interest rate parity’, *The Journal of Finance* **73**(3), 915–957.
- El-Mekkaoui, M. & Flood, M. D. (1998), ‘Put-call parity revisited: Intraday evidence from the options market’, *Journal of International Financial Markets, Institutions and Money* **8**(3–4), 357–376.
- European Securities and Markets Authority (2024), ‘Markets in crypto-assets regulation (mica)’, Regulatory information page. MiCA became applicable in stages in 2024. Accessed 2026-06-04.
URL: <https://www.esma.europa.eu/esmas-activities/digital-finance-and-innovation/markets-crypto-assets-regulation-mica>
- Felföldi-Szűcs, N. (2024), ‘Put-call parity in crypto option markets’, *Finance Research Letters* **61**, 104963.
- Gârleanu, N. & Pedersen, L. H. (2011), ‘Margin-based asset pricing and deviations from the law of one price’, *The Review of Financial Studies* **24**(6), 1980–2022.

- Kamara, A. & Miller, T. W. (1995), ‘Daily and intradaily tests of european put-call parity’, *The Journal of Financial and Quantitative Analysis* **30**(4), 519–539.
- Klemkosky, R. C. & Resnick, B. G. (1979), ‘Put-call parity and market efficiency’, *The Journal of Finance* **34**(5), 1141–1155.
- Králik, B. (2026a), ‘publications: Amu-liquidity-premium’, GitHub repository. Accessed 2026-07-16.
URL: <https://github.com/balazs14/publications/tree/main/AMU-liquidity-premium>
- Králik, B. (2026b), ‘Tardisproject’, GitHub repository. Accessed 2026-07-16.
URL: <https://github.com/balazs14/TardisProject>
- Lucic, V. & Sepp, A. (2024), ‘Valuation and hedging of cryptocurrency inverse options’, *Quantitative Finance* .
- Nisbet, M. (1992), ‘Put-call parity theory and an empirical test of the efficiency of the london traded options market’, *Journal of Banking & Finance* **16**(2), 381–403.
- Pan, K. & Zeng, Y. (2017), Etf arbitrage under liquidity mismatch, Working Paper 59, European Systemic Risk Board.
- Shleifer, A. & Vishny, R. W. (1997), ‘The limits of arbitrage’, *The Journal of Finance* **52**(1), 35–55.
- Tardis.dev (2026), ‘Tardis: Historical crypto market data api’, Product and documentation website. Accessed 2026-06-16.
URL: <https://tardis.dev/>
- Todorov, K. (2021), ‘The anatomy of bond etf arbitrage’, *BIS Quarterly Review* pp. 41–53.
- U.S. Securities and Exchange Commission (2005), ‘Regulation nms, release no. 34-51808: Order protection (trade-through) rule’, Final Rule, 17 CFR Parts 200, 201, 230, 240, 242, 249, and 270. Accessed 2026-07-16.
URL: <https://www.sec.gov/rules/final/34-51808.pdf>
- U.S. Securities and Exchange Commission (2024a), ‘Order granting accelerated approval of proposed rule changes to list and trade shares of ether-based exchange-traded products’, SEC Release, May 2024. Accessed 2026-06-04.
URL: <https://www.sec.gov/>
- U.S. Securities and Exchange Commission (2024b), ‘Statement on the approval of spot bitcoin exchange-traded products’, SEC Chair Statement, January 10, 2024. Accessed 2026-06-04.
URL: <https://www.sec.gov/newsroom/speeches-statements/gensler-statement-spot-bitcoin-011023>